

# Three tools for high-throughput computing with GAMESS

Riccardo Murri

`<riccardo.murri@uzh.ch>`

Grid Computing Competence Centre, University of Zurich

<http://www.gc3.uzh.ch/>

May 5, 2011



University of Zurich

# Three tools for high-throughput GAMESS

**GRunDB** Compute energy of a set of molecules.

**GGamess** Run an arbitrary set of GAMESS .INP files in parallel.

**gc3libs.template** Generate a set of files from a given template.

The purpose of this talk is to determine whether these can be of any use to you or GC3 should stop supporting them.

# What is GRunDB then?

GRunDB is a tool to automate:

- running GAMESS on a data-bank (GMTKN24) of molecule geometries
- comparing computed stoichiometry results with known-good ones

## Functional high-level view

GRunDB is implemented as a Linux command-line tool on top of the GC3Libs/GC3Utils toolkit.

GRunDB takes as input:

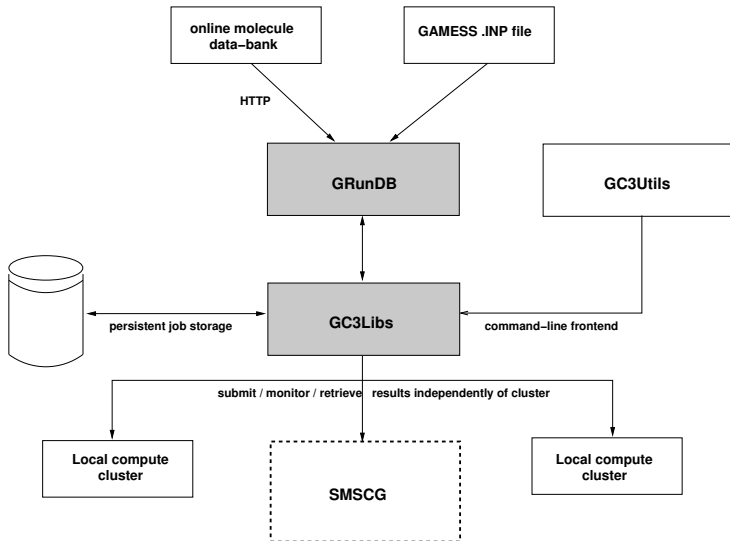
- a set of molecules (a subset of an online data-bank: either the original GMTKN24 or its local copy at UZH)
- a GAMESS input file (except the \$DATA section)

It creates a GAMESS job for each molecule in the subset, plugging its geometry in the template input file.

When all jobs are done:

- results are extracted from the output files
- a summary table compares computed energy with reference data

# Architecture



# The GAMESS.UZH molecule database: Subset index

University of Zurich > Organic Chemistry Institute

Baldridge Research Group | Contact | Sitemap

University of Zurich

## The GAMESS.UZH molecule database

[Edit](#) | [RecentChanges](#) | [Preferences](#) | [Discussion](#)

### Welcome to the GAMESS.UZH molecule database

The available molecules are divided into subsets; the list of all subsets is available below. For each molecule, a table of reference stoichiometry data is given, accompanied by the molecule geometry in GAMESS .inp file format.

This site is organized as a [Wiki](#); only authorized users can [edit pages](#) to add or modify content, but anyone can view the pages and download data files.

The database is maintained by people in the Baldridge group at the [Organic Chemistry Institute](#) of the [University of Zürich](#); IT support is provided by the [Grid Computing Competence Centre](#).

### Available molecule sets

The available molecules are divided into subsets; the list of all subsets is available below. For each molecule, a table of reference stoichiometry data is given, accompanied by the molecule geometry in downloadable GAMESS .inp file format.

- [ACONF](#)
- [AL2X](#)
- [BH76](#)
- [BH76RC](#)
- [BHPERI](#)
- [CYCONF](#)
- [DARC](#)
- [DC9](#)
- [G21EA](#)
- [G21IP](#)
- [G2RC](#)
- [IDISP](#)
- [ISO34](#)
- [MB08-165](#)
- [NBRC](#)
- [O3ADD6](#)
- [PA](#)
- [PCONE](#)

# The GAMESS.UZH molecule database: The AL2X subset

University of Zurich > Organic Chemistry Institute > The GAMESS.UZH molecule database > subsets

Beldridge Research Group | Contact | Sitemap

The GAMESS.UZH molecule database

[Edit](#) | [RecentChanges](#) | [Preferences](#) | [Discussion](#)

## The AL2X subset

Original data taken from: Johnson, E. R.; Mori-Sanchez, P.; Cohen, A. J.; Yang, W. *J. Chem. Phys.*, 2008, 129, 204112. (back-corrected experimental values); all values are in kcal/mol.

### Reference data

For each reaction, the relevant systems' names, the stoichiometry and the reference value are given. The systems' names refer to the geometry files (see section "GAMESS input" below). Negative values in the stoichiometry columns refer to reactants, positive values to products.

#	Systems	Stoichiometry	Ref.
0	alh3 al2h6	2 -1	35.80
1	al2f6 alf3	-1 2	52.60
2	al2cl6 alcl3	-1 2	31.40
3	albr3 al2br6	2 -1	28.40
4	alme2 al2me4	2 -1	37.80
5	alme2 alme3 al2me5	1 1 -1	30.00
6	alme3 al2me6	2 -1	21.40

[Direct data download](#)

### Reference GAMESS input

- [al2br6.inp](#)
- [al2cl6.inp](#)
- [al2f6.inp](#)
- [al2h6.inp](#)
- [al2me4.inp](#)
- [al2me5.inp](#)
- [al2me6.inp](#)
- [albr3.inp](#)
- [alcl3.inp](#)
- [alf3.inp](#)

# Start analysis of a molecule set

GRunDB creates a job for each molecule in the specified subset. A *session* name must be given to record the current analysis.

Subsets can be specified by their name on the web page. More than one (or *ALL*) can be analyzed in a single GRunDB go.

```
$ ./grundb.py new AL2X session1
```

Input file name	State (JobID)	Info
al2br6	NEW (job.680)	New at Mon Nov 29 14:06:46 2010
al2cl6	NEW (job.681)	New at Mon Nov 29 14:06:46 2010
al2f6	NEW (job.682)	New at Mon Nov 29 14:06:46 2010
al2h6	NEW (job.683)	New at Mon Nov 29 14:06:46 2010
...		
alf3	NEW (job.689)	New at Mon Nov 29 14:06:47 2010
alh3	NEW (job.690)	New at Mon Nov 29 14:06:47 2010
alme2	NEW (job.691)	New at Mon Nov 29 14:06:47 2010
alme3	NEW (job.692)	New at Mon Nov 29 14:06:47 2010



# Submit jobs to the Grid

Once a session has been created, a single command invocation is needed to submit jobs to SMSCG or University clusters.

```
$ ./grundb.py progress session1
```

```
Insert AAI/Switch password for user m1058036 :
```

```
Queue selected: all.q@idgc3grid01.uzh.ch
```

```
File uploaded: /tmp/rmurri/rsl.q0p5wn
```

```
File uploaded: /home/rmurri/gc3/gc3utils/0.10/grundb/take1.inp.d/AL2X/al2f6
```

```
...
```

Input file name	State (JobID)	Info
al2br6	SUBMITTED (job.680)	Submitted at Mon Nov 29 14:08:04 2010
al2cl6	SUBMITTED (job.681)	Submitted at Mon Nov 29 14:07:53 2010
al2f6	SUBMITTED (job.682)	Submitted at Mon Nov 29 14:07:29 2010
al2h6	SUBMITTED (job.683)	Submitted at Mon Nov 29 14:07:41 2010
...		
alh3	SUBMITTED (job.690)	Submitted at Mon Nov 29 14:07:50 2010
alme2	SUBMITTED (job.691)	Submitted at Mon Nov 29 14:07:59 2010
alme3	SUBMITTED (job.692)	Submitted at Mon Nov 29 14:08:02 2010

# Monitor job progress and execution

The same command is used to monitor job execution.

```
$ ./grundb.py progress session1
```

Input file name	State (JobID)	Info
al2br6	SUBMITTED (job.680)	Submitted at Mon Nov 29 14:08:04 2010
al2cl6	RUNNING (job.681)	Running at Mon Nov 29 14:08:40 2010
al2f6	RUNNING (job.682)	Running at Mon Nov 29 14:08:40 2010
al2h6	RUNNING (job.683)	Running at Mon Nov 29 14:08:40 2010
...		
alh3	RUNNING (job.690)	Running at Mon Nov 29 14:08:40 2010
alme2	SUBMITTED (job.691)	Submitted at Mon Nov 29 14:07:59 2010
alme3	SUBMITTED (job.692)	Submitted at Mon Nov 29 14:08:02 2010

## Output retrieval and post-processing

Again, grunldb progress will automatically retrieve and post-process results of jobs that have finished execution, extracting the energy values needed to compute stoichiometry results.

```
$ ./grunldb.py progress session1
```

```
File downloaded: gsiftp://idgc3grid01.uzh.ch:2811/jobs/17057129103608390550
```

```
File downloaded: gsiftp://idgc3grid01.uzh.ch:2811/jobs/17057129103608390550
```

```
...
```

Input file name	State (JobID)	Info
al2br6	RUNNING (job.680)	Running at Mon Nov 29 14:09:00 2010
al2cl6	RUNNING (job.681)	Running at Mon Nov 29 14:08:40 2010
al2f6	DONE (job.682)	Final-b2plyp energy= -1084.1929109480
al2h6	DONE (job.683)	Final-b2plyp energy= -488.2225951188
...		
alh3	DONE (job.690)	Final-b2plyp energy= -244.0835095562
alme2	DONE (job.691)	Final-b2plyp energy= -322.6947308201
alme3	DONE (job.692)	Final-b2plyp energy= -361.9996423212

# Finally...

When all jobs are done, GRunDB computes stoichiometry data and compares it to the reference data.

```
$ ./grundb.py progress session1
```

Input file name	State (JobID)	Info
al2br6	DONE (job.682)	Final r-m06 energy is -15929.9575541891
...		after 19 iterations
alme3	DONE (job.694)	Final r-m06 energy is -362.1226427375
		after 18 iterations

## STOICHIOMETRY DATA

Reaction	Comp. energy	(Ref. data; deviation)
=====		
	AL2X	
-----		
2*alh3 + -1*al2h6	+35.70	(+35.80; -0.10)
-1*al2f6 + 2*alf3	+48.46	(+52.60; -4.14)
-1*al2cl6 + 2*alcl3	+28.15	(+31.40; -3.25)
2*albr3 + -1*al2br6	+25.41	(+28.40; -2.99)
2*alme2 + -1*al2me4	+34.52	(+37.80; -3.28)
1*alme2 + 1*alme3 + -1*al2me5	+28.55	(+30.00; -1.45)
2*alme3 + -1*al2me6	+21.72	(+21.40; +0.32)

# What is GGameSS then?

GGameSS is a command-line tool to submit a set of GAMESS .INP files in parallel.

- For each job, manage the entire lifecycle: submit, monitor, retrieve output when done.
- Each job is independent of others.
- You can stop and restart it at a later time, processing continues from where it was interrupted.
- Finally exits when all jobs are done.

# Example: running the GAMESS tests / 1

Running ggameSS once submits the the jobs.

```
$ ls tests
exam01.inp  exam02.inp  exam03.inp  exam04.inp  exam05.inp
[...]
exam41.inp  exam42.inp  exam43.inp  exam44.inp

$ ./ggameSS.py -r ocikbpra tests/
Status of jobs in the 'ggameSS' session: (at 11:14:23, 05/05/11)
      NEW      0/44      (0.0%)
      STOPPED  0/44      (0.0%)
      SUBMITTED 44/44     (100.0%)
      TERMINATED 0/44      (0.0%)
      TERMINATING 0/44      (0.0%)
      total    44/44     (100.0%)
```

## Example: running the GAMESS tests / 2

Running it again updates status and fetches results of finished jobs.

You can also request a detailed listing of the jobs with the `-l` option:

```
$ ./ggameSS.py -l -r ocikbptra tests/
```

JobID	Job name	State	Info
job.8945	exam42	SUBMITTED	Submitted to 'smscg' at Thu May 5 11:11:58 2011
job.8944	exam16	SUBMITTED	Submitted to 'smscg' at Thu May 5 11:12:02 2011
[...]			
job.8937	exam20	TERMINATED	Execution of gamess terminated normally thu may
job.8936	exam29	TERMINATED	Execution of gamess terminated normally thu may

Each job has its own output directory.

```
$ ls exam29
```

```
exam29.dat exam29.out
```

## Example: running the GAMESS tests / 3

Perhaps the most interesting thing is that you can tell ggameSS to keep running until all jobs are done and their output retrieved.

Example: keep running, and update job status every 20 seconds.

```
$ ./ggameSS.py -r ocikbpra tests/ -C 20
Status of jobs in the 'ggameSS' session: (at 11:27:48, 05/05/11)
      NEW      0/44      (0.0%)
      STOPPED   0/44      (0.0%)
      SUBMITTED 5/44      (11.4%)
      TERMINATED 39/44     (88.6%)
      TERMINATING 0/44     (0.0%)
      ok        39/44     (88.6%)
      total     44/44     (100.0%)
...continues running
```



# What is gc3libs.template?

gc3libs.template is a tool for generating a set of files from a template.

It's a *programming library*, not a command-line tool; you need to do some Python language programming to exploit it fully.

## Example: Timm's "GAMESS benchmark" / 1

To use `gc3libs.template` you need:

- ① A string with the template contents of the file;
- ② actual parameters to be substituted in the template;
- ③ an optional "filter function" to determine which combination of parameters are acceptable.

```
GAMESS_INP = Template("""
  $$CONTRL RUNTYP=ENERGY MAXIT=1 UNITS=BOHR $$END
  $$CONTRL $SCF ISPHER=$ISPHER $$END
  $$ACCURACY ITOL=$ITOL ILOAD=$ILOAD $$END
  $$SYSTEM MWORDS=10 $$END
  $$BASIS $BASIS $$END
  $$GUESS GUESS=HUCKEL $$END

  $$DATA
$GEOMETRY
  $$END
""",
    match_ispher_with_basis,
```

## Example: Timm's "GAMESS benchmark" / 2

Substitution parameters can be simple list of values (including multi-line strings).

```

GEOMETRY = [
    """Water
C1
O   8.0      0.0      0.0      0.0
H   1.0      0.0      1.428036  1.0957706
H   1.0      0.0     -1.428036  1.0957706""",
    """Methane
Td
C   6.0    0.0      0.0      0.0
H   1.0    0.6252197764  0.6252197764  0.6252197764""",
],

```

## Example: Timm's "GAMESS benchmark" / 3

But they can be other templates, which are recursively expanded.

```

BASIS = [Template("GBASIS=$SIMPLEBASIS",
    SIMPLEBASIS = ["MINI", "MIDI", "DZV", "TZV",
        "CCD", "CCT", "CCQ", "CC5", "CC6"]),
    Template("GBASIS=$GBASIS NGAUSS=$NGAUSS $NPFUNC $NDFUNC $NFFUNC",
        acceptable_gbasis_and_ngauss,
        GBASIS = ["STO", "N21", "N31", "N311"],
        NGAUSS = [2, 3, 4, 5, 6],
        NPFUNC = ["", "NPFUNC=1"],
        NDFUNC = ["", "NDFUNC=1"],
        NFFUNC = ["", "NFFUNC=1"],
    ),
],

```

Thank you!

(Any questions?)

# Three tools for high-throughput GAMESS

**GRunDB** Compute energy of a set of molecules.

**GGamess** Run an arbitrary set of GAMESS .INP files in parallel.

**gc3libs.template** Generate a set of files from a given template.

## Well, and the errors...?

Errors happen, and it's not necessarily a users' fault. (e.g., disk full on a remote cluster)

How does GRunDB deal with this?

## Well, and the errors...?

Errors happen, and it's not necessarily a users' fault. (e.g., disk full on a remote cluster)

How does GRunDB deal with this?

Use GC3Utils to get control over the *individual* job:

- Inspect job output and logs in the *session.out.d* directory (or use the `gtail` command)
- Take measures against the failure
- Re-submit the failed job with `gresub`
- GRunDB will notice and re-process the results when the new job is done.